

# A4\_Final

February 3, 2023

## 1 Data Preprocessing

```
[282]: import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np
import os
import pathlib
import random
```

```
[283]: path = '/Users/richardreynard/Downloads/AY 2022/SP53:22/MA3832/Assignment4/
↳Dataset/'
data_dir = pathlib.Path(path)
```

```
[284]: class_names = np.array([sorted(item.name for item in data_dir.glob("*"))])
class_names
```


```
[284]: array([[ 'DS_Store', 'Mild_Demented', 'Moderate_Demented',
           'Non_Demented', 'Very_Mild_Demented']], dtype='<U18')
```

```
[285]: imageCount = len(list(data_dir.glob("*/*.jpg")))
imageCount
```

```
[285]: 6400
```


```
[286]: import cv2
from matplotlib import pyplot as plt
non_demented_image = cv2.imread("Dataset/Non_Demented/non.jpg")
plt.imshow(non_demented_image)
```

```
[286]: <matplotlib.image.AxesImage at 0x7fd257fb1c40>
```




```
[287]: very_mild_demented_image = cv2.imread("Dataset/Very_Mild_Demented/verymild.jpg")  
plt.imshow(very_mild_demented_image)
```

```
[287]: <matplotlib.image.AxesImage at 0x7fd1da92fc10>
```




```
[288]: mild_demented_image = cv2.imread("Dataset/Mild_Demented/mild.jpg")  
plt.imshow(mild_demented_image)
```

```
[288]: <matplotlib.image.AxesImage at 0x7fd2b9eb8700>
```



```
[289]: moderate_demented_image = cv2.imread("Dataset/Moderate_Demented/moderate.jpg")  
plt.imshow(moderate_demented_image)
```

```
[289]: <matplotlib.image.AxesImage at 0x7fce41252700>
```



```
[296]: batch_size = 32  
img_height = 256  
img_width = 256
```

```
[297]: from tensorflow.keras.utils import image_dataset_from_directory  
from tensorflow.keras.utils import image_dataset_from_directory  
  
train_data = image_dataset_from_directory(  
    data_dir,  
    validation_split=0.2,  
    subset="training",  
    seed=123,  
    image_size=(img_height, img_width),  
    batch_size=batch_size)  
  
val_data = image_dataset_from_directory(data_dir,  
                                       validation_split=0.2,  
                                       subset="validation",  
                                       seed=123,  
                                       image_size=(img_height, img_width),  
                                       batch_size=batch_size)
```

Found 6400 files belonging to 4 classes.  
Using 5120 files for training.  
Found 6400 files belonging to 4 classes.  
Using 1280 files for validation.

## 2 Build Model

```
[302]: from tensorflow.keras import layers

model = tf.keras.Sequential([

    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),

    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),

    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),

    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),

    layers.Conv2D(128, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),

    layers.Dropout(0.5),
    layers.Flatten(),

    layers.Dense(128, activation='relu'),
    layers.Dense(4, activation="softmax")
])
```

```
[303]: model.compile(optimizer="Adam",
                    loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                    metrics=["accuracy"])
model.summary()
```

Model: "sequential\_61"

Layer (type)	Output Shape	Param #
rescaling_32 (Rescaling)	(None, 256, 256, 3)	0
conv2d_106 (Conv2D)	(None, 256, 256, 16)	448
max_pooling2d_121 (MaxPooling2D)	(None, 128, 128, 16)	0

conv2d_107 (Conv2D)	(None, 128, 128, 32)	4640
max_pooling2d_122 (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_108 (Conv2D)	(None, 64, 64, 64)	18496
max_pooling2d_123 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_109 (Conv2D)	(None, 32, 32, 128)	73856
max_pooling2d_124 (MaxPooling2D)	(None, 16, 16, 128)	0
dropout_44 (Dropout)	(None, 16, 16, 128)	0
flatten_44 (Flatten)	(None, 32768)	0
dense_125 (Dense)	(None, 128)	4194432
dense_126 (Dense)	(None, 4)	516

```

=====
Total params: 4,292,388
Trainable params: 4,292,388
Non-trainable params: 0
-----

```

```

[304]: epochs = 10
history = model.fit(train_data,
                    epochs=epochs,
                    validation_data=val_data,
                    batch_size=batch_size)

```

```

Epoch 1/10
160/160 [=====] - 220s 1s/step - loss: 1.0449 -
accuracy: 0.4977 - val_loss: 0.9187 - val_accuracy: 0.5703
Epoch 2/10
160/160 [=====] - 206s 1s/step - loss: 0.9055 -
accuracy: 0.5680 - val_loss: 0.8020 - val_accuracy: 0.6367
Epoch 3/10
160/160 [=====] - 222s 1s/step - loss: 0.7925 -
accuracy: 0.6359 - val_loss: 0.7002 - val_accuracy: 0.7133
Epoch 4/10
160/160 [=====] - 213s 1s/step - loss: 0.6673 -
accuracy: 0.7014 - val_loss: 0.5533 - val_accuracy: 0.7719

```

```

Epoch 5/10
160/160 [=====] - 204s 1s/step - loss: 0.5319 -
accuracy: 0.7719 - val_loss: 0.4236 - val_accuracy: 0.8203
Epoch 6/10
160/160 [=====] - 200s 1s/step - loss: 0.3935 -
accuracy: 0.8393 - val_loss: 0.2939 - val_accuracy: 0.8875
Epoch 7/10
160/160 [=====] - 231s 1s/step - loss: 0.2847 -
accuracy: 0.8871 - val_loss: 0.2009 - val_accuracy: 0.9336
Epoch 8/10
160/160 [=====] - 232s 1s/step - loss: 0.2060 -
accuracy: 0.9207 - val_loss: 0.1634 - val_accuracy: 0.9438
Epoch 9/10
160/160 [=====] - 234s 1s/step - loss: 0.1659 -
accuracy: 0.9377 - val_loss: 0.1267 - val_accuracy: 0.9641
Epoch 10/10
160/160 [=====] - 231s 1s/step - loss: 0.1204 -
accuracy: 0.9502 - val_loss: 0.1228 - val_accuracy: 0.9547

```

```

[312]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(32,32))
plt.subplot(4,3,4)
plt.plot(epochs_range,acc,label='Accuracy')
plt.plot(epochs_range,val_acc,label="Validation Accuracy")
plt.plot(epochs_range,loss,label='Loss')
plt.plot(epochs_range,val_loss,label="Validation Loss")
plt.legend()


```

```

[312]: <matplotlib.legend.Legend at 0x7fcf999b5520>

```





### 3 Transfer Learning

```
[313]: import tensorflow as tf
from keras_preprocessing.image import ImageDataGenerator
import os
import cv2
import random
from matplotlib import pyplot as plt
import pathlib
import pandas as pd
import numpy as np
```

```
[239]: # train_dir = '/Users/richardreynard/Downloads/Dataset3/Train_Set'
# test_dir = '/Users/richardreynard/Downloads/Dataset3/Test_Set'
```

```
[314]: train_dir = '/Users/richardreynard/Downloads/Dataset_BiT/Training Data'
test_dir = '/Users/richardreynard/Downloads/Dataset_BiT/Testing Data'
val_dir = '/Users/richardreynard/Downloads/Dataset_BiT/Validation Data'
```

```
[238]: # from keras.preprocessing.image import ImageDataGenerator as IDG

# # Initialize image data generator
# train_gen = IDG(rescale=1/255, rotation_range=10, horizontal_flip=True,
    ↪vertical_flip=False)
# test_gen = IDG(rescale=1/255, rotation_range=10, horizontal_flip=True,
    ↪vertical_flip=False)
# valid_gen = IDG(rescale=1/255)

# # Load the datasets
# train_ds = train_gen.flow_from_directory(train_dir, shuffle=True,
    ↪batch_size=64, target_size=(256,256), class_mode='binary')
# test_ds = test_gen.flow_from_directory(train_dir, shuffle=True,
    ↪batch_size=64, target_size=(256,256), class_mode='binary')
# valid_ds = test_gen.flow_from_directory(test_dir, shuffle=True,
    ↪batch_size=32, target_size=(256,256), class_mode='binary')
```

Found 5120 images belonging to 4 classes.

Found 5120 images belonging to 4 classes.

Found 1280 images belonging to 4 classes.

```
[240]: from keras.preprocessing.image import ImageDataGenerator as IDG

# Initialize image data generator
train_gen = IDG(rescale=1/255, rotation_range=10, horizontal_flip=True,
    ↪vertical_flip=False)
test_gen = IDG(rescale=1/255, rotation_range=10, horizontal_flip=True,
    ↪vertical_flip=False)
valid_gen = IDG(rescale=1/255)

# Load the datasets
train_ds = train_gen.flow_from_directory(train_dir, shuffle=True,
    ↪batch_size=64, target_size=(256,256), class_mode='binary')
test_ds = test_gen.flow_from_directory(test_dir, shuffle=True, batch_size=64,
    ↪target_size=(256,256), class_mode='binary')
valid_ds = test_gen.flow_from_directory(val_dir, shuffle=True, batch_size=32,
    ↪target_size=(256,256), class_mode='binary')
```

Found 4267 images belonging to 4 classes.

Found 1422 images belonging to 4 classes.

Found 711 images belonging to 4 classes.

```
[241]: from keras.layers import Dense, GlobalAveragePooling2D as GAP, Dropout
from keras.models import load_model, Sequential

# Pre Trained Models
```

```
from tensorflow.keras.applications import ResNet50V2, InceptionV3, Xception, \
↳ ResNet50, ResNet152V2
```

### 3.1 Using ResNet50V2 - Most Accurate

```
[242]: # Base Model
base = ResNet50V2(include_top=False, input_shape=(256,256,3))
base.trainable = False

# Model Architecture
model = tf.keras.Sequential([
    base,
    GAP(),
    Dense(1024, kernel_initializer='he_normal', activation='relu'),
    Dense(512, kernel_initializer='he_normal', activation='relu'),
    Dropout(0.4),
    Dense(4, activation="softmax"),
])

# Compile
model.compile(
    loss='sparse_categorical_crossentropy',
    optimizer=tf.keras.optimizers.Adam(learning_rate=2e-3),
    metrics=['accuracy']
)
```

```
[243]: epochs = 20
history = model.fit(
    train_ds,
    validation_data = valid_ds,
    epochs = epochs)
```

```
Epoch 1/20
67/67 [=====] - 399s 6s/step - loss: 1.5325 - accuracy:
0.4870 - val_loss: 0.9243 - val_accuracy: 0.5682
Epoch 2/20
67/67 [=====] - 413s 6s/step - loss: 0.9041 - accuracy:
0.5742 - val_loss: 0.8858 - val_accuracy: 0.5851
Epoch 3/20
67/67 [=====] - 426s 6s/step - loss: 0.8566 - accuracy:
0.5948 - val_loss: 0.8882 - val_accuracy: 0.5767
Epoch 4/20
67/67 [=====] - 492s 7s/step - loss: 0.8359 - accuracy:
0.6072 - val_loss: 0.8668 - val_accuracy: 0.5977
Epoch 5/20
67/67 [=====] - 620s 9s/step - loss: 0.8000 - accuracy:
```

0.6307 - val\_loss: 0.8594 - val\_accuracy: 0.5668  
Epoch 6/20  
67/67 [=====] - 469s 7s/step - loss: 0.7669 - accuracy:  
0.6445 - val\_loss: 0.8376 - val\_accuracy: 0.6020  
Epoch 7/20  
67/67 [=====] - 456s 7s/step - loss: 0.7536 - accuracy:  
0.6478 - val\_loss: 0.8160 - val\_accuracy: 0.6188  
Epoch 8/20  
67/67 [=====] - 491s 7s/step - loss: 0.7149 - accuracy:  
0.6773 - val\_loss: 0.8063 - val\_accuracy: 0.6357  
Epoch 9/20  
67/67 [=====] - 508s 8s/step - loss: 0.6793 - accuracy:  
0.7003 - val\_loss: 0.8001 - val\_accuracy: 0.6414  
Epoch 10/20  
67/67 [=====] - 497s 7s/step - loss: 0.6428 - accuracy:  
0.7209 - val\_loss: 0.8043 - val\_accuracy: 0.6371  
Epoch 11/20  
67/67 [=====] - 501s 7s/step - loss: 0.6205 - accuracy:  
0.7270 - val\_loss: 0.7921 - val\_accuracy: 0.6385  
Epoch 12/20  
67/67 [=====] - 513s 8s/step - loss: 0.6232 - accuracy:  
0.7288 - val\_loss: 0.7934 - val\_accuracy: 0.6371  
Epoch 13/20  
67/67 [=====] - 503s 8s/step - loss: 0.6022 - accuracy:  
0.7413 - val\_loss: 0.8621 - val\_accuracy: 0.6343  
Epoch 14/20  
67/67 [=====] - 513s 8s/step - loss: 0.5771 - accuracy:  
0.7570 - val\_loss: 0.8078 - val\_accuracy: 0.6442  
Epoch 15/20  
67/67 [=====] - 490s 7s/step - loss: 0.5540 - accuracy:  
0.7628 - val\_loss: 0.8008 - val\_accuracy: 0.6582  
Epoch 16/20  
67/67 [=====] - 437s 7s/step - loss: 0.5424 - accuracy:  
0.7654 - val\_loss: 0.7863 - val\_accuracy: 0.6428  
Epoch 17/20  
67/67 [=====] - 481s 7s/step - loss: 0.5419 - accuracy:  
0.7628 - val\_loss: 0.8347 - val\_accuracy: 0.6245  
Epoch 18/20  
67/67 [=====] - 487s 7s/step - loss: 0.4900 - accuracy:  
0.7968 - val\_loss: 0.8192 - val\_accuracy: 0.6498  
Epoch 19/20  
67/67 [=====] - 496s 7s/step - loss: 0.4868 - accuracy:  
0.7935 - val\_loss: 0.8406 - val\_accuracy: 0.6568  
Epoch 20/20  
67/67 [=====] - 478s 7s/step - loss: 0.4421 - accuracy:  
0.8188 - val\_loss: 0.7834 - val\_accuracy: 0.6793

```
[244]: epochs = 50
        history = model.fit(
            train_ds,
            validation_data = valid_ds,
            epochs = epochs)
```

Epoch 1/50

67/67 [=====] - 437s 7s/step - loss: 0.4423 - accuracy: 0.8186 - val\_loss: 0.8957 - val\_accuracy: 0.6245

Epoch 2/50

67/67 [=====] - 464s 7s/step - loss: 0.4435 - accuracy: 0.8106 - val\_loss: 0.8037 - val\_accuracy: 0.6667

Epoch 3/50

67/67 [=====] - 464s 7s/step - loss: 0.4367 - accuracy: 0.8275 - val\_loss: 0.9119 - val\_accuracy: 0.6118

Epoch 4/50

67/67 [=====] - 445s 7s/step - loss: 0.4227 - accuracy: 0.8245 - val\_loss: 0.8668 - val\_accuracy: 0.6062

Epoch 5/50

67/67 [=====] - 442s 7s/step - loss: 0.4400 - accuracy: 0.8172 - val\_loss: 0.9114 - val\_accuracy: 0.6146

Epoch 6/50

67/67 [=====] - 445s 7s/step - loss: 0.4066 - accuracy: 0.8294 - val\_loss: 0.8820 - val\_accuracy: 0.6610

Epoch 7/50

67/67 [=====] - 455s 7s/step - loss: 0.3823 - accuracy: 0.8423 - val\_loss: 0.8471 - val\_accuracy: 0.6624

Epoch 8/50

67/67 [=====] - 470s 7s/step - loss: 0.4095 - accuracy: 0.8306 - val\_loss: 0.8988 - val\_accuracy: 0.6287

Epoch 9/50

67/67 [=====] - 478s 7s/step - loss: 0.3966 - accuracy: 0.8434 - val\_loss: 0.7246 - val\_accuracy: 0.6920

Epoch 10/50

67/67 [=====] - 481s 7s/step - loss: 0.3556 - accuracy: 0.8570 - val\_loss: 0.8094 - val\_accuracy: 0.6554

Epoch 11/50

67/67 [=====] - 482s 7s/step - loss: 0.3332 - accuracy: 0.8688 - val\_loss: 0.8156 - val\_accuracy: 0.6737

Epoch 12/50

67/67 [=====] - 498s 7s/step - loss: 0.4041 - accuracy: 0.8310 - val\_loss: 0.8227 - val\_accuracy: 0.6751

Epoch 13/50

67/67 [=====] - 484s 7s/step - loss: 0.3727 - accuracy: 0.8449 - val\_loss: 0.7421 - val\_accuracy: 0.7032

Epoch 14/50

67/67 [=====] - 513s 8s/step - loss: 0.3273 - accuracy:

0.8683 - val\_loss: 1.0109 - val\_accuracy: 0.6470  
 Epoch 15/50  
 67/67 [=====] - 546s 8s/step - loss: 0.3265 - accuracy:  
 0.8688 - val\_loss: 0.9632 - val\_accuracy: 0.6512  
 Epoch 16/50  
 67/67 [=====] - 484s 7s/step - loss: 0.3104 - accuracy:  
 0.8774 - val\_loss: 0.9600 - val\_accuracy: 0.6470  
 Epoch 17/50  
 67/67 [=====] - 469s 7s/step - loss: 0.3155 - accuracy:  
 0.8723 - val\_loss: 0.8573 - val\_accuracy: 0.6821  
 Epoch 18/50  
 67/67 [=====] - 464s 7s/step - loss: 0.2857 - accuracy:  
 0.8835 - val\_loss: 0.7902 - val\_accuracy: 0.7032  
 Epoch 19/50  
 67/67 [=====] - 477s 7s/step - loss: 0.3040 - accuracy:  
 0.8777 - val\_loss: 0.8240 - val\_accuracy: 0.6737  
 Epoch 20/50  
 67/67 [=====] - 464s 7s/step - loss: 0.3072 - accuracy:  
 0.8795 - val\_loss: 1.1248 - val\_accuracy: 0.6315  
 Epoch 21/50  
 67/67 [=====] - 453s 7s/step - loss: 0.2597 - accuracy:  
 0.8999 - val\_loss: 0.9245 - val\_accuracy: 0.6962  
 Epoch 22/50  
 67/67 [=====] - 449s 7s/step - loss: 0.2850 - accuracy:  
 0.8875 - val\_loss: 0.9493 - val\_accuracy: 0.7060  
 Epoch 23/50  
 67/67 [=====] - 433s 6s/step - loss: 0.3066 - accuracy:  
 0.8749 - val\_loss: 0.8603 - val\_accuracy: 0.7060  
 Epoch 24/50  
 67/67 [=====] - 437s 7s/step - loss: 0.2648 - accuracy:  
 0.8983 - val\_loss: 0.9196 - val\_accuracy: 0.6807  
 Epoch 25/50  
 67/67 [=====] - 445s 7s/step - loss: 0.2656 - accuracy:  
 0.9013 - val\_loss: 1.0210 - val\_accuracy: 0.6610  
 Epoch 26/50  
 67/67 [=====] - 461s 7s/step - loss: 0.2500 - accuracy:  
 0.9049 - val\_loss: 0.8302 - val\_accuracy: 0.6920  
 Epoch 27/50  
 67/67 [=====] - 432s 6s/step - loss: 0.2547 - accuracy:  
 0.8938 - val\_loss: 0.8668 - val\_accuracy: 0.6850  
 Epoch 28/50  
 67/67 [=====] - 426s 6s/step - loss: 0.2568 - accuracy:  
 0.8964 - val\_loss: 0.9620 - val\_accuracy: 0.6681  
 Epoch 29/50  
 67/67 [=====] - 435s 6s/step - loss: 0.2518 - accuracy:  
 0.9041 - val\_loss: 0.8521 - val\_accuracy: 0.7018  
 Epoch 30/50  
 67/67 [=====] - 462s 7s/step - loss: 0.2481 - accuracy:

0.8995 - val\_loss: 0.9633 - val\_accuracy: 0.6709  
Epoch 31/50  
67/67 [=====] - 458s 7s/step - loss: 0.2333 - accuracy:  
0.9058 - val\_loss: 0.8234 - val\_accuracy: 0.7046  
Epoch 32/50  
67/67 [=====] - 461s 7s/step - loss: 0.2401 - accuracy:  
0.9077 - val\_loss: 0.8987 - val\_accuracy: 0.7103  
Epoch 33/50  
67/67 [=====] - 465s 7s/step - loss: 0.2328 - accuracy:  
0.9063 - val\_loss: 0.8510 - val\_accuracy: 0.7004  
Epoch 34/50  
67/67 [=====] - 451s 7s/step - loss: 0.2246 - accuracy:  
0.9105 - val\_loss: 0.9158 - val\_accuracy: 0.7089  
Epoch 35/50  
67/67 [=====] - 446s 7s/step - loss: 0.2418 - accuracy:  
0.9070 - val\_loss: 1.0032 - val\_accuracy: 0.6667  
Epoch 36/50  
67/67 [=====] - 451s 7s/step - loss: 0.2316 - accuracy:  
0.9067 - val\_loss: 1.0431 - val\_accuracy: 0.6850  
Epoch 37/50  
67/67 [=====] - 447s 7s/step - loss: 0.2166 - accuracy:  
0.9138 - val\_loss: 0.9070 - val\_accuracy: 0.7328  
Epoch 38/50  
67/67 [=====] - 452s 7s/step - loss: 0.2319 - accuracy:  
0.9131 - val\_loss: 0.9205 - val\_accuracy: 0.7229  
Epoch 39/50  
67/67 [=====] - 457s 7s/step - loss: 0.2477 - accuracy:  
0.9041 - val\_loss: 0.9027 - val\_accuracy: 0.7018  
Epoch 40/50  
67/67 [=====] - 447s 7s/step - loss: 0.2351 - accuracy:  
0.9070 - val\_loss: 0.8493 - val\_accuracy: 0.7089  
Epoch 41/50  
67/67 [=====] - 443s 7s/step - loss: 0.1828 - accuracy:  
0.9320 - val\_loss: 1.1205 - val\_accuracy: 0.6934  
Epoch 42/50  
67/67 [=====] - 453s 7s/step - loss: 0.1863 - accuracy:  
0.9313 - val\_loss: 1.0778 - val\_accuracy: 0.6878  
Epoch 43/50  
67/67 [=====] - 462s 7s/step - loss: 0.1897 - accuracy:  
0.9266 - val\_loss: 1.2172 - val\_accuracy: 0.6835  
Epoch 44/50  
67/67 [=====] - 444s 7s/step - loss: 0.1833 - accuracy:  
0.9288 - val\_loss: 1.0887 - val\_accuracy: 0.6765  
Epoch 45/50  
67/67 [=====] - 446s 7s/step - loss: 0.2024 - accuracy:  
0.9222 - val\_loss: 0.8976 - val\_accuracy: 0.7173  
Epoch 46/50  
67/67 [=====] - 461s 7s/step - loss: 0.1749 - accuracy:

```
0.9356 - val_loss: 0.8974 - val_accuracy: 0.7145
Epoch 47/50
67/67 [=====] - 456s 7s/step - loss: 0.1829 - accuracy:
0.9285 - val_loss: 0.9759 - val_accuracy: 0.7060
Epoch 48/50
67/67 [=====] - 444s 7s/step - loss: 0.2186 - accuracy:
0.9135 - val_loss: 1.0988 - val_accuracy: 0.7201
Epoch 49/50
67/67 [=====] - 444s 7s/step - loss: 0.1880 - accuracy:
0.9262 - val_loss: 1.0354 - val_accuracy: 0.6934
Epoch 50/50
67/67 [=====] - 449s 7s/step - loss: 0.2136 - accuracy:
0.9161 - val_loss: 1.0307 - val_accuracy: 0.6835
```

```
[246]: acc = history.history['accuracy']
      val_acc = history.history['val_accuracy']


      loss = history.history['loss']
      val_loss = history.history['val_loss']

      epochs_range = range(epochs)

      plt.figure(figsize=(8,8))
      plt.subplot(1,2,2)
      plt.plot(epochs_range,acc,label='Accuracy')
      plt.plot(epochs_range,val_acc,label="Validation Accuracy")
      plt.plot(epochs_range,loss,label='Loss')
      plt.plot(epochs_range,val_loss,label="Validation Loss")
      plt.legend()
```

```
[246]: <matplotlib.legend.Legend at 0x7fd27e9af730>
```





### 3.2 Using BiT - Big Transfer

```
[222]: # import os
# import numpy as np
# import pandas as pd
# from glob import glob
# import tensorflow as tf

# ! pip install "tensorflow>=2.0.0"
```

```

# ! pip install --upgrade tensorflow-hub
# import tensorflow_hub as hub
# from IPython.display import clear_output as cls

# # Data
# from tensorflow.keras.utils import load_img, img_to_array
# from keras.preprocessing.image import ImageDataGenerator

# # Data Visualization
# import plotly.express as px
# import matplotlib.pyplot as plt

# # Model
# from keras.models import Sequential, load_model
# from keras.layers import GlobalAvgPool2D as GAP, Dense, Dropout

# # Callbacks
# from keras.callbacks import EarlyStopping, ModelCheckpoint

# # Pre-Trained Model
# from tensorflow.keras.applications import ResNet50V2

```

```

[224]: # Check Training Data Information
train_path = '/Users/richardreynard/Downloads/Dataset_BiT/Training Data/'
class_names = sorted(os.listdir(train_path))
n_classes = len(class_names)

# Show
print(f"Total Number of Classes : {n_classes} \nClass Names : {class_names}")

```

```

Total Number of Classes : 5
Class Names : ['.DS_Store', 'Mild_Demented', 'Moderate_Demented',
'Non_Demented', 'Very_Mild_Demented']

```

```

[225]: # # Check Testing Data Information
test_path = '/Users/richardreynard/Downloads/Dataset_BiT/Testing Data/'
class_names = sorted(os.listdir(test_path))
n_classes = len(class_names)

# Show
print(f"Total Number of Classes : {n_classes} \nClass Names : {class_names}")

```

```

Total Number of Classes : 5
Class Names : ['.DS_Store', 'Mild_Demented', 'Moderate_Demented',
'Non_Demented', 'Very_Mild_Demented']

```

```
[227]: # Check Validation Data Information
valid_path = '/Users/richardreynard/Downloads/Dataset_BiT/Validation Data/'
class_names = sorted(os.listdir(valid_path))
n_classes = len(class_names)

# Show
print(f"Total Number of Classes : {n_classes} \nClass Names : {class_names}")
```

Total Number of Classes : 5  
Class Names : ['.DS\_Store', 'Mild\_Demented', 'Moderate\_Demented',  
'Non\_Demented', 'Very\_Mild\_Demented']

```
[228]: # Initialize Generator
train_gen = ImageDataGenerator(rescale=1/255., rotation_range=10,
    ↪horizontal_flip=True)
valid_gen = ImageDataGenerator(rescale=1/255.)
test_gen = ImageDataGenerator(rescale=1/255)

# Load Data
train_ds = train_gen.flow_from_directory(train_path, class_mode='binary',
    ↪target_size=(256,256), shuffle=True, batch_size=32)
valid_ds = valid_gen.flow_from_directory(valid_path, class_mode='binary',
    ↪target_size=(256,256), shuffle=True, batch_size=32)
test_ds = test_gen.flow_from_directory(test_path, class_mode='binary',
    ↪target_size=(256,256), shuffle=True, batch_size=32)
```

Found 4267 images belonging to 4 classes.  
Found 711 images belonging to 4 classes.  
Found 1422 images belonging to 4 classes.

```
[229]: # Import BiT model
bit_model_url = "https://tfhub.dev/google/bit/m-r50x1/1"
bit_module = hub.KerasLayer(bit_model_url)
```

```
[230]: model = Sequential([
    bit_module,
    Dense(4, activation='softmax', kernel_initializer='zeros')
], name='bit-custom')
```

```
[231]: BATCH_SIZE = 32
lr = 1e-3 * BATCH_SIZE/512
print(f"Learning rate : {lr}")
```

Learning rate : 6.25e-05

```
[232]: SCHEDULE_BOUNDARIES = [
    200,
```

```
300,  
400,  
]
```

```
[233]: lr_schedule = tf.keras.optimizers.schedules.PiecewiseConstantDecay(  
        boundaries=SCHEDULE_BOUNDARIES,  
        values=[  
            lr,  
            lr * 0.1,  
            lr * 0.01,  
            lr * 0.001,  
        ],  
    )  
optimizer = tf.keras.optimizers.SGD(learning_rate=lr_schedule, momentum=0.9)
```

```
[234]: model.compile(  
        loss='sparse_categorical_crossentropy',  
        optimizer=optimizer,  
        metrics=['accuracy']  
    )
```

```
[235]: history = model.fit(train_ds, validation_data=valid_ds, epochs=4)
```

```
Epoch 1/4  
134/134 [=====] - 993s 7s/step - loss: 1.0017 -  
accuracy: 0.5245 - val_loss: 0.9224 - val_accuracy: 0.5668  
Epoch 2/4  
134/134 [=====] - 936s 7s/step - loss: 0.9179 -  
accuracy: 0.5664 - val_loss: 0.8909 - val_accuracy: 0.5696  
Epoch 3/4  
134/134 [=====] - 923s 7s/step - loss: 0.8731 -  
accuracy: 0.5814 - val_loss: 0.8920 - val_accuracy: 0.5823  
Epoch 4/4  
134/134 [=====] - 949s 7s/step - loss: 0.8693 -  
accuracy: 0.5857 - val_loss: 0.8917 - val_accuracy: 0.5809
```

### 3.3 Using EfficientNetB0 - Least Efficient

```
[315]: train_dir = '/Users/richardreynard/Downloads/Dataset_BiT/Training Data'  
test_dir = '/Users/richardreynard/Downloads/Dataset_BiT/Testing Data'  
val_dir = '/Users/richardreynard/Downloads/Dataset_BiT/Validation Data'
```

```
[316]: # Initialize Generator  
train_gen = ImageDataGenerator(rescale=1/255., rotation_range=10,   
    ↪ horizontal_flip=True)  
valid_gen = ImageDataGenerator(rescale=1/255.)  
test_gen = ImageDataGenerator(rescale=1/255)
```

```

# Load Data
train_ds = train_gen.flow_from_directory(train_path, class_mode='binary',
    ↳target_size=(256,256), shuffle=True, batch_size=32)
valid_ds = valid_gen.flow_from_directory(valid_path, class_mode='binary',
    ↳target_size=(256,256), shuffle=True, batch_size=32)
test_ds = test_gen.flow_from_directory(test_path, class_mode='binary',
    ↳target_size=(256,256), shuffle=True, batch_size=32)

```

Found 4267 images belonging to 4 classes.

Found 711 images belonging to 4 classes.

Found 1422 images belonging to 4 classes.

```

[317]: base2 = tf.keras.applications.EfficientNetB0(include_top=False,
    ↳input_shape=(256,256,3))
base2.trainable = False
model2 = tf.keras.Sequential([
    base2,
    GAP(),
    Dense(1024, kernel_initializer='he_normal', activation='relu'),
    Dropout(0.4),
    Dense(512, kernel_initializer='he_normal', activation='relu'),
    Dropout(0.4),
    Dense(4, activation="softmax")
])

# Compile
model2.compile(
    loss='sparse_categorical_crossentropy',
    optimizer=tf.keras.optimizers.Adam(learning_rate=2e-3),
    metrics=['accuracy']
)

```

```

[318]: epochs = 5
history = model2.fit(
    train_ds,
    validation_data = valid_ds,
    epochs = epochs)

```

Epoch 1/5

134/134 [=====] - 258s 2s/step - loss: 1.1474 -

accuracy: 0.4596 - val\_loss: 1.0611 - val\_accuracy: 0.5007

Epoch 2/5

134/134 [=====] - 234s 2s/step - loss: 1.0635 -

accuracy: 0.4821 - val\_loss: 1.0434 - val\_accuracy: 0.5007

Epoch 3/5

```
134/134 [=====] - 266s 2s/step - loss: 1.0633 -  
accuracy: 0.4872 - val_loss: 1.0414 - val_accuracy: 0.5007  
Epoch 4/5  
134/134 [=====] - 264s 2s/step - loss: 1.0460 -  
accuracy: 0.4896 - val_loss: 1.0349 - val_accuracy: 0.5007  
Epoch 5/5  
134/134 [=====] - 245s 2s/step - loss: 1.0376 -  
accuracy: 0.5008 - val_loss: 1.0391 - val_accuracy: 0.5007
```

```
[326]: acc = history.history['accuracy']  
val_acc = history.history['val_accuracy']  
loss = history.history['loss']  
val_loss = history.history['val_loss']  
  
epochs_range = range(epochs)  
  
plt.figure(figsize=(10,8))  
plt.subplot(2,2,2)  
plt.plot(epochs_range,acc,label='Accuracy')  
plt.plot(epochs_range,val_acc,label="Validation Accuracy")  
plt.plot(epochs_range,loss,label='Loss')  
plt.plot(epochs_range,val_loss,label="Validation Loss")  
plt.legend()
```

```
[326]: <matplotlib.legend.Legend at 0x7fd07539a730>
```

